

Reflection — Building My Portfolio with Generative AI

李冠穎 Kuan-Yin Lee · Student ID: TBD · HW1, 2026 Generative AI

Process. I approached this assignment the way I approach system design at work: requirements first. Before asking the AI for anything, I wrote my constraints in a small file (RQ.md): publish only the website, never expose the raw resume files in my `info/` folder, keep a versioned log of every question and change, and write commit messages that someone could audit later. I then asked Claude to build the whole assignment from those requirements and my existing resume and autobiography. The AI extracted my experience, papers and patents, read the course handout, proposed a static single-page site with a GitHub Actions workflow that deploys only the `site/` folder, and produced a first version in one pass. It then rendered the page on a desktop and a phone viewport and corrected layout bugs it found in its own screenshots.

Human edits and judgment. The most valuable part of the collaboration was deciding what *not* to accept. My resume contains more than a public website should: phone number, the context of a PhD application, and a rejected paper submission. I kept only e-mail and city, removed the application framing, and turned the rejected submission into a short note about learning from reviewer feedback without naming the venue. When the AI had no source for my co-authors' Chinese names, it kept their English names rather than guessing — a rule I want in every AI-assisted document. My two resumes disagree on my employment dates; instead of silently choosing one, the conflict was recorded as an open question. I also asked that the architecture diagram be labelled "illustrative", because a simplified picture should not be mistaken for the figure in my paper.

What I learned about human–AI collaboration. First, a written specification makes the AI far more useful than a clever prompt: it keeps both of us anchored and makes the output checkable. Second, the AI is fast at breadth — layout, bilingual content, dark mode, deployment scripts — but the judgment about privacy, honesty and emphasis stayed with me. Third, momentum and clarification pull against each other. When the AI stopped to ask questions, I preferred that it make a reasonable assumption and log it; the QA table in SPEC.md turned that into a to-do list rather than a blocker. Finally, traceability matters. Because every change is recorded in HISTORY.md and linked from commit messages, I can explain exactly what the AI generated, why I accepted or changed it, and how the result improved — which is precisely what this assignment asks of me, and what my own research on multi-agent systems requires of autonomous agents.